

Computer Science 384
St. George Campus

Monday, January 20, 2020
University of Toronto

Homework Assignment #1: Search
Due: Tuesday, February 4, 2020 by 10:00 PM

Silent Policy: A silent policy will take effect 24 hours before this assignment is due, i.e. no question about this assignment will be answered, whether it is asked on the discussion board, via email or in person.

Late Policy: 10% per day after the use of 3 grace days.

Total Marks: This assignment represents 11% of the course grade.

Handing in this Assignment

What to hand in on paper: Nothing.

What to hand in electronically: You must submit your assignment electronically. Download `solution.py`, `sokoban.py`, `search.py`, `tips.txt`, `autograder.py` and `acknowledgment_form.pdf` from: <http://www.teach.cs.utoronto.ca/~csc384h/winter/Assignments/A1/>. These are in the zip file that is called `search.zip`. Modify `solution.py` so that it solves the Sokoban problem as specified in this document. Then, submit your modified `solution.py` and `tips.txt` and your signed `acknowledgment_form.pdf` using MarkUs. Your login to MarkUs is your teach.cs username and password. You can submit a new version of any file at any time, though the lateness penalty applies if you submit after the deadline. For the purposes of determining the lateness penalty, the submission time is considered to be the time of your latest submission.

We will test your code electronically. You will be supplied with a testing script that will run a **subset** of the tests. If your code fails all of the tests performed by the script (using **Python version 3.7**), you will receive a failing grade on the assignment.

When your code is submitted, we will run a more extensive set of tests which will include the tests run in the provided testing script and a number of other tests. You have to pass all of these more elaborate tests to obtain full marks on the assignment.

Your code will not be evaluated for partial correctness; it either works or it doesn't. It is your responsibility to hand in something that passes at least some of the tests in the provided testing script.

- *Make certain that your code runs on teach.cs using python3 (version 3.7) using only standard imports.* This version is installed as “python3” on teach.cs. Your code will be tested using this version and you will receive zero marks if it does not run using this version.
- *Do not add any non-standard imports from within the python file you submit (the imports that are already in the template files must remain).* Once again, non-standard imports will cause your code to fail the testing and you will receive zero marks.
- *Do not change the supplied starter code.* Your code will be tested using the original starter code, and if it relies on changes you made to the starter code, you will receive zero marks.

Clarification Page: Important corrections (hopefully few or none) and clarifications to the assignment will be posted on the Assignment 1 Clarification page: http://www.teach.cs.toronto.edu/~csc384h/winter/Assignments/A1/a1_fa.html.

You are responsible for monitoring the A1 Clarification page.

Help Sessions: There will be two help sessions for this assignment. Dates and times for these sessions will be posted to the course website and to Piazza ASAP.

Questions: Questions about the assignment should be asked on Piazza:




Figure 1: A state of the Sokoban puzzle.

<https://piazza.com/utoronto.ca/winter2020/csc384/home>.

If you have a question of a personal nature, please email one of the A1 TAs, Randy (rhickey at cs.toronto.edu) or Sam (sam.motamed at mail.utoronto.ca) or a course instructor. Make sure to place [CSC384] and A1 in the subject line of your message.

1 Introduction

The goal of this assignment will be to implement a working solver for the puzzle game Sokoban shown in Figure 1. Sokoban is a puzzle game in which a warehouse robot must push boxes into storage spaces. The rules hold that only one box can be moved at a time, that boxes can only be pushed by a robot and not pulled, and that neither robots nor boxes can pass through obstacles (walls or other boxes). In addition, robots cannot push more than one box, i.e., if there are two boxes in a row, the robot cannot push them. The game is over when all the boxes are in their storage spots.

In our version of Sokoban the rules are slightly more complicated, as there may be more than one warehouse robot available to push boxes. These robots cannot pass through one another nor can they move simultaneously, however.

Sokoban can be played online at <https://www.sokobanonline.com/play>. We recommend that you familiarize yourself with the rules and objective of the game before proceeding, although it is worth noting that the version that is presented online is only an example. We will give a formal description of the puzzle in the next section

2 Description of Sokoban

Sokoban has the following formal description. Read the description carefully. Note that our version differs from the standard one.

- The puzzle is played on a rectangle board that is a grid *board* with N squares in the x -dimension and M squares in the y -dimension.
- Each state contains the x and y coordinates for each robot, the boxes, the storage spots, and the obstacles.
- From each state, each robot can move North, South, East, or West. No two robots can move simultaneously, however. If a robot moves to the location of a box, the box will move one square in the same direction. Boxes and robots cannot pass through walls or obstacles, however. Robots cannot push more than one box at a time; if two boxes are in succession the robot will not be able to move them. Movements that cause a box to move more than one unit of the grid are also illegal. Whether or not a robot is pushing an object does not change the cost.
- Each movement is of equal cost.
- The goal is achieved when each box is located in a storage area on the grid.

Ideally, we will want our robots to organize everything before the supervisor arrives. This means that with each problem instance, you will be given a computation time constraint. You must attempt to provide some legal solution to the problem (i.e. a plan) within this constraint. Better plans will be plans that are shorter, i.e. that require fewer operators to complete.

Your goal is to implement anytime algorithms for this problem: ones that generates better solutions (i.e. shorter plans) the more computation time they are given.

3 Code You Have Been Provided

The file `search.py`, which is available from the website, provides a generic search engine framework and code to perform several different search routines. This code will serve as a base for your Sokoban solver. A brief description of the functionality of `search.py` follows. The code itself is documented and worth reading.

- An object of class `StateSpace` represents a node in the state space of a generic search problem. The base class defines a fixed interface that is used by the `SearchEngine` class to perform search in that state space.

For the Sokoban problem, we will define a concrete sub-class that inherits from `StateSpace`. This concrete sub-class will inherit some of the “utility” methods that are implemented in the base class. Each `StateSpace` object s has the following key attributes:

- $s.gval$: the g value of that node, i.e., the cost of getting to that state.
- $s.parent$: the parent `StateSpace` object of s , i.e., the `StateSpace` object that has s as a successor. This will be `None` if s is the initial state.
- $s.action$: a string that contains that name of the action that was applied to $s.parent$ to generate s . Will be “`START`” if s is the initial state.

- An object of class `SearchEngine` *se* runs the search procedure. A `SearchEngine` object is initialized with a search strategy (`'depth_first'`, `'breadth_first'`, `'best_first'`, `'a_star'`, or `'custom'`) and a cycle checking level (`'none'`, `'path'`, or `'full'`).

Note that `SearchEngine` depends on two auxiliary classes:

- An object of class `sNode` *sn* which represents a node in the search space. Each object *sn* contains a `StateSpace` object and additional details: *hval*, i.e., the heuristic function value of that state and *gval*, i.e. the cost to arrive at that node from the initial state. An *fval_fn* and *weight* are tied to search nodes during the execution of a search, where applicable.
- An object of class `Open` is used to represent the search frontier. The search frontier will be organized in the way that is appropriate for a given search strategy.

When a `SearchEngine`'s search strategy is set to `'custom'`, you will have to specify the way that *f* values of nodes are calculated; these values will structure the order of the nodes that are expanded during your search.

Once a `SearchEngine` object has been instantiated, you can set up a specific search with:

```
init_search(initial_state, goal_fn, heur_fn, fval_fn)
```

and execute that search with

```
search(timebound, costbound)
```

The arguments are as follows:

- *initial_state* will be an object of type `StateSpace`; it is your start state.
- *goal_fn(s)* is a function which returns `True` if a given state *s* is a goal state and `False` otherwise.
- *heuristic_fn(s)* is a function that returns a heuristic value for state *s*. This function will only be used if your search engine has been instantiated to be a heuristic search (e.g. *best_first*).
- *fval_fn(sNode, weight)* defines *f* values for states. This function will only be used by your search engine if it has been instantiated to execute a `'custom'` search. Note that this function takes in an *sNode* and that an *sNode* contains not only a state but additional measures of the state (e.g. a *gval*). The function also takes in a float *weight*. It will use the variables that are provided to arrive at an *f* value calculation for the state contained in the *sNode*.
- *timebound* is a bound on the amount of time your code will be allowed to execute the search. Once the run time exceeds the time bound, the search must stop; if no solution has been found, the search will return *False*.
- *costbound* is an optional parameter that is used to set boundaries on the cost of nodes that are explored. This *costbound* is defined as a list of three values. *costbound*[0] is used to prune states based on their *g*-values; any state with a *g*-value higher than *costbound*[0] will not be expanded. *costbound*[1] is used to prune states based on their *h*-values; any state with an *h*-value higher than *costbound*[1] will not be expanded. Finally, *costbound*[2] is used to prune states based on their *f*-values; any state with an *f*-value higher than *costbound*[2] will not be expanded.

For this assignment we have also provided `sokoban.py`, which specializes `StateSpace` for the Sokoban problem. You will therefore not need to encode representations of Sokoban states or the successor function for Sokoban! These have been provided to you so that you can focus on implementing good search heuristics and anytime algorithms.

The file `sokoban.py` contains:

- An object of class `sokobanState`, which is a `StateSpace` with these additional key attributes:
 - `s.width`: the width of the Sokoban board
 - `s.height`: the height of the Sokoban board
 - `s.robots`: positions for each robot that is on the board. Each robot position is a tuple (x,y) , that denotes the robot's x and y position.
 - `s.bboxes`: positions for each box that is on the board. Each box position is also an (x,y) tuple.
 - `s.storage`: positions for each storage bin that is on the board (also (x,y) tuples).
 - `s.obstacles`: positions for obstacles that are on the board. Obstacles, like robots and boxes, are also tuples of (x,y) coordinates.
- `sokobanState` also contains the following key functions:
 - `successors()`: This function generates a list of `SokobanStates` that are successors to a given `SokobanState`. Each state will be annotated by the action that was used to arrive at the `SokobanState`. These actions are (r,d) tuples wherein r denotes the index of the robot that moved d denotes the direction of movement of the robot.
 - `hashable_state()`: This is a function that calculates a unique index to represents a particular `sokobanState`. It is used to facilitate path and cycle checking.
 - `print_state()`: This function prints a `sokobanState` to stdout.

Note that `sokobanState` depends on one auxiliary class:

- An object of class `Direction`, which is used to define the directions that each robot can move and the effect of this movement.

Also note that `sokoban.py` contains a set of 20 initial states for Sokoban problems, which are stored in the tuple `PROBLEMS`. You can use these states to test your implementations.

The file `solution.py` contains the methods that need to be implemented.

The file `autograder.py` runs some tests on your code to give you an indication of how well your methods perform.

4 Assignment Specifics

To complete this assignment you must modify `solution.py` to:

- Implement a **Manhattan distance heuristic** (*heur_manhattan_distance(state)*). This heuristic will be used to estimate how many moves a current state is from a goal state. Your implementation should calculate the **sum of Manhattan distances between each box that has yet to be stored and the storage point nearest to it**. Ignore the positions of obstacles in your calculations and assume that many boxes can be stored at one location.
- Implement **Anytime Greedy Best-First Search** (*anytime_gbfs(initial_state;heur_fn;timebound)*). Details regarding this algorithm are provided in Section 5.
- Implement **Anytime Weighted A*** (*weighted_astar(initial_state,timebound)*). Details about this algorithm are provided in Section 6. Note that your implementation will require you to instantiate a `SearchEngine` object with a custom search strategy. To do this you must therefore define an *f*-value function (*fval_function(sNode,weight)*) and remember to provide this when you execute *init_search*.
- Implement a non-trivial heuristic for Sokoban that improves on the Manhattan distance heuristic (*heur_alternate(state)*).
- You should give **three tips (2 sentences each)** as if you were advising someone who was attempting this problem for this first time on what to do. Write these tips in the file `tips.txt`.

Note that when we are testing your code, we will limit each run of your algorithm on `teach.cs` to a fixed time bound. Instances that are not solved within this limit will provide an interesting evaluation metric: failure rate.

5 Anytime Greedy Best-First Search

Greedy best-first search **expands nodes with lowest $h(node)$ first**. The solution found by this algorithm may not be optimal. Anytime greedy best-first search (which is called *anytime_gbfs* in the code) continues searching after a solution is found in order to improve solution quality. Since we have found a path to the goal after the first iteration, we can introduce a cost bound for pruning: if node has $g(node)$ greater than the best path to the goal found so far, we can prune it. The algorithm returns either when we have expanded all non-pruned nodes, in which case the best solution found by the algorithm is the optimal solution, or when it runs out of time. We **prune based on the g_value** of the node only because greedy best-first search is not necessarily run with an admissible heuristic.

Record the time when *anytime_gbfs* is called with `os.times()[0]`. Each time you call search, you should update the time bound with the remaining allowed time. The automarking script will confirm that your algorithm obeys the specified time bound.

6 Anytime Weighted A*

Instead of A*'s regular node-valuation formula ($f = g(node) + h(node)$), Weighted A* introduces a weighted formula:

$$f = g(node) + w * h(node)$$

where $g(\text{node})$ is the cost of the path to node , $h(\text{node})$ is the estimated cost of getting from node to the goal, and $w > 1$ is a bias towards states that are closer to the goal. Theoretically, the smaller w is, the better the solution will be (i.e. the closer to the optimal solution it will be ... *why and under what assumptions??*). However, different values of w will require different computation times.

Weighted A* typically starts with a value for w that is more than 1. Since the first solution that is found by Weighted A* may not be optimal if $w > 1$, we can keep searching after we have found a solution.

Anytime Weighted A* continues to search, decreasing w at each iteration, until either there are no nodes left to expand (and our best solution is the optimal one) or it runs out of time. Since a solution we find in one iteration may not be optimal, we can introduce a cost bound for pruning during subsequent iterations: if node has a $g(\text{node}) + h(\text{node})$ value that is greater than the cost of the best path to the goal found so far, we can prune it.

When you are passing in an f_val function to $init_search$ for this problem, you will need to have specified the weight for the $fval_function$. You can do this by wrapping the $fval_function(sN, weight)$ you have written in an anonymous function, i.e.,

$$wrapped_fval_function = (\text{lambda } sN : fval_function(sN, weight))$$



Figure 2: A state of the Water Jugs puzzle.

7 StateSpace **Example:** WaterJugs.py

WaterJugs.py contains an example implementation of the search engine for the Water Jugs problem shown in Figure 2.

- You have two containers that can be used to store water. One has a three-gallon capacity, and the other has a four-gallon capacity. Each has an initial, known, amount of water in it.
- You have the following actions available:
 - You can fill either container from the tap until it is full.
 - You can dump the water from either container.
 - You can pour the water from one container into the other, until either the source container is empty or the destination container is full.
- You are given a goal amount of water to have in each container. You are trying to achieve that goal in the minimum number of actions, assuming the actions have uniform cost.

WaterJugs.py has an implementation of the Water Jugs puzzle that is suitable for using with search.py. Note that in addition to implementing the three key methods of StateSpace, the author has created a set of tests that show how to operate the search engine. You should study these to see how the search engine works.

GOOD LUCK!